

Register Allocation

Xiao Jia

April 25th, 2012

Outline

- Introduction
- Graph coloring
- Linear scan

Introduction

- Want to replace temporary variables with some fixed set of registers

Introduction

- Want to replace temporary variables with some fixed set of registers
- We will judge this phase by ...
 - Code review
 - Limiting # of executed instructions, e.g. no greater than 1 million
 - *Compare: only 160,000 for my optimized 8-queen*

Graph coloring

- **First:** need to know which variables are live after each instruction
 - Two simultaneously live variables cannot be allocated to the same register

Graph coloring

Control Flow Graph

- For every node **n** in CFG, we have **out[n]**
 - Set of temporaries live out of n
- Two variables *interfere* if
 - both initially live (i.e. function arguments), or
 - both appear in out[n] for any n, or
 - one is defined and the other is in out[n]
 - $x = b - c$ where x is dead & b is live interfere
- How to assign registers to variables?

Interference graph

- **Nodes** of the graph = variables
- **Edges** connect variables that interfere with one another
- Nodes will be assigned a **color** corresponding to the register assigned to the variable
- Two colors can't be next to one another in the graph

Interference graph

Instructions	Live vars
--------------	-----------

$b = a + 2$	
-------------	--

$c = b * b$	
-------------	--

$b = c + 1$	
-------------	--

return $b * a$	
----------------	--

Interference graph

Instructions

Live vars

$b = a + 2$

$c = b * b$

$b = c + 1$

return $b * a$

b, a

Interference graph

Instructions	Live vars
--------------	-----------

$b = a + 2$	
-------------	--

$c = b * b$	
-------------	--

	a, c
--	--------

$b = c + 1$	
-------------	--

	b, a
--	--------

$\text{return } b * a$	
------------------------	--

Interference graph

Instructions	Live vars
--------------	-----------

$b = a + 2$	
-------------	--

	b, a
--	--------

$c = b * b$	
-------------	--

	a, c
--	--------

$b = c + 1$	
-------------	--

	b, a
--	--------

$\text{return } b * a$	
------------------------	--

Interference graph

Instructions	Live vars
	a
b = a + 2	
	b,a
c = b * b	
	a,c
b = c + 1	
	b,a
return b * a	

Interference graph

Instructions

$b = a + 2$

$c = b * b$

$b = c + 1$

return $b * a$

Live vars


a

a, b

a, c

a, b

color	register
	\$t1
	\$t2



Interference graph

Instructions

$b = a + 2$

$c = b * b$

$b = c + 1$

return $b * a$

Live vars


a

a, b

a, c

a, b

color	register
	\$t1
	\$t2



Graph coloring

- Questions:
 - Can we efficiently find a coloring of the graph whenever possible?
 - Can we efficiently find the optimum coloring of the graph?
 - What do we do when there aren't enough colors (registers) to color the graph?

Coloring a graph


- Kempe's algorithm [1879] for finding a K -coloring of a graph
- **Step 1 (simplify):** find a node with **at most $K-1$** edges and cut it out of the graph.
(Remember this node on a stack for later stages.)

Coloring a graph

- Once a coloring is found for the simpler graph, we can always color the node we saved on the stack
- **Step 2 (color):** when the simplified subgraph has been colored, add back the node on the top of the stack and assign it a color not taken by one of the adjacent nodes

Coloring


color	register
	\$t1
	\$t2



stack:

Coloring

color	register
	\$t1
	\$t2




stack:

c

Coloring

color	register
	\$t1
	\$t2




stack:

e
c

Coloring

color	register
	\$t1
	\$t2




stack:

a
e
c

Coloring

color	register
	\$t1
	\$t2




stack:

b
a
e
c

Coloring

color	register
	\$t1
	\$t2




stack:

d
b
a
e
c

Coloring

color	register
	\$t1
	\$t2




stack:

b
a
e
c

Coloring

color	register
	\$t1
	\$t2




stack:

a
e
c

Coloring

color	register
	\$t1
	\$t2




stack:

e
c

Coloring

color	register
	\$t1
	\$t2




stack:

c

Coloring

color	register
	\$t1
	\$t2




stack:

Failure

- If the graph cannot be colored, it will eventually be simplified to graph in which **every node has at least K neighbors**
- Sometimes, the graph is still K -colorable!
- Finding a K -coloring in all situations is an **NP-complete** problem
 - We will have to approximate to make register allocators fast enough

Coloring


color	register
	\$t1
	\$t2



stack:

Coloring

color	register
	\$t1
	\$t2




stack:
d

all nodes have
2 neighbours!

Coloring

color	register
	\$t1
	\$t2




stack:

b
d

Coloring

color	register
	\$t1
	\$t2




stack:

c
e
a
b
d

Coloring

color	register
	\$t1
	\$t2




stack:

e
a
b
d

Coloring

color	register
	\$t1
	\$t2




stack:

a
b
d

Coloring

color	register
	\$t1
	\$t2




stack:

b
d

Coloring

color	register
	\$t1
	\$t2




stack:

d

Coloring

color	register
	\$t1
	\$t2




stack:

We got lucky!

Coloring

color	register
	\$t1
	\$t2

Some graphs can't be colored
in K colors:




stack:

c
b
e
a
d

Coloring

color	register
	\$t1
	\$t2

Some graphs can't be colored
in K colors:




stack:

b
e
a
d

Coloring

color	register
	\$t1
	\$t2

Some graphs can't be colored
in K colors:




stack:

e
a
d

Coloring

color	register
	\$t1
	\$t2

Some graphs can't be colored
in K colors:



no colors left for e!

stack:


e
a
d

Spilling

- **Step 3 (spilling):** once all nodes have K or more neighbors, pick a node for **spilling**
 - Storage on the stack
- There are many heuristics that can be used to pick a node
 - not in an inner loop

Spilling code

- We need to generate extra instructions to load variables from stack and store them
- These instructions use registers themselves. What to do?
 - **Stupid approach:** always keep extra registers handy for shuffling data in and out: **what a waste!**
 - **Better approach:** ?



Dedicated registers

Spilling code

- We need to generate extra instructions to load variables from stack and store them
- These instructions use registers themselves. What to do?
 - **Stupid approach:** always keep extra registers handy for shuffling data in and out: **what a waste!**
 - **Better approach:** rewrite code introducing a new temporary; rerun liveness analysis and register allocation

Rewriting code

- Consider: `add t1, t1, t2`
 - Suppose `t2` is selected for spilling and assigned to stack location `24($fp)`
 - Introduce new temporary `t3` for just this instruction and rewrite:
 - `ld t3, 24($fp)`
 - `add t1, t1, t3`
 - Advantage: `t3` has a very short live range and is much less likely to interfere.
 - Rerun the algorithm; fewer variables will spill

Precolored Nodes

- Some variables are pre-assigned to registers
 - Frame pointer
 - Arguments (\$a0, \$a1, \$a2, \$a3)
 - Function call defines (trashies) caller-save registers
- Treat these registers as special temporaries; before beginning, **add them to the graph with their colors**

Precolored Nodes

See Tiger
book for more
details


- Can't simplify a graph by removing a precolored node
- Precolored nodes are the starting point of the coloring process
- Once simplified down to colored nodes start adding back the other nodes as before

Optimizing Moves

- Code generation produces a lot of extra move instructions
 - `mov t1, t2`
 - If we can assign t1 and t2 to the same register, we do not have to execute the `mov`
 - Idea: if t1 and t2 are not connected in the interference graph, we **coalesce** into a single variable

Coalescing

- Problem: coalescing can increase the number of interference edges and make a graph uncolorable

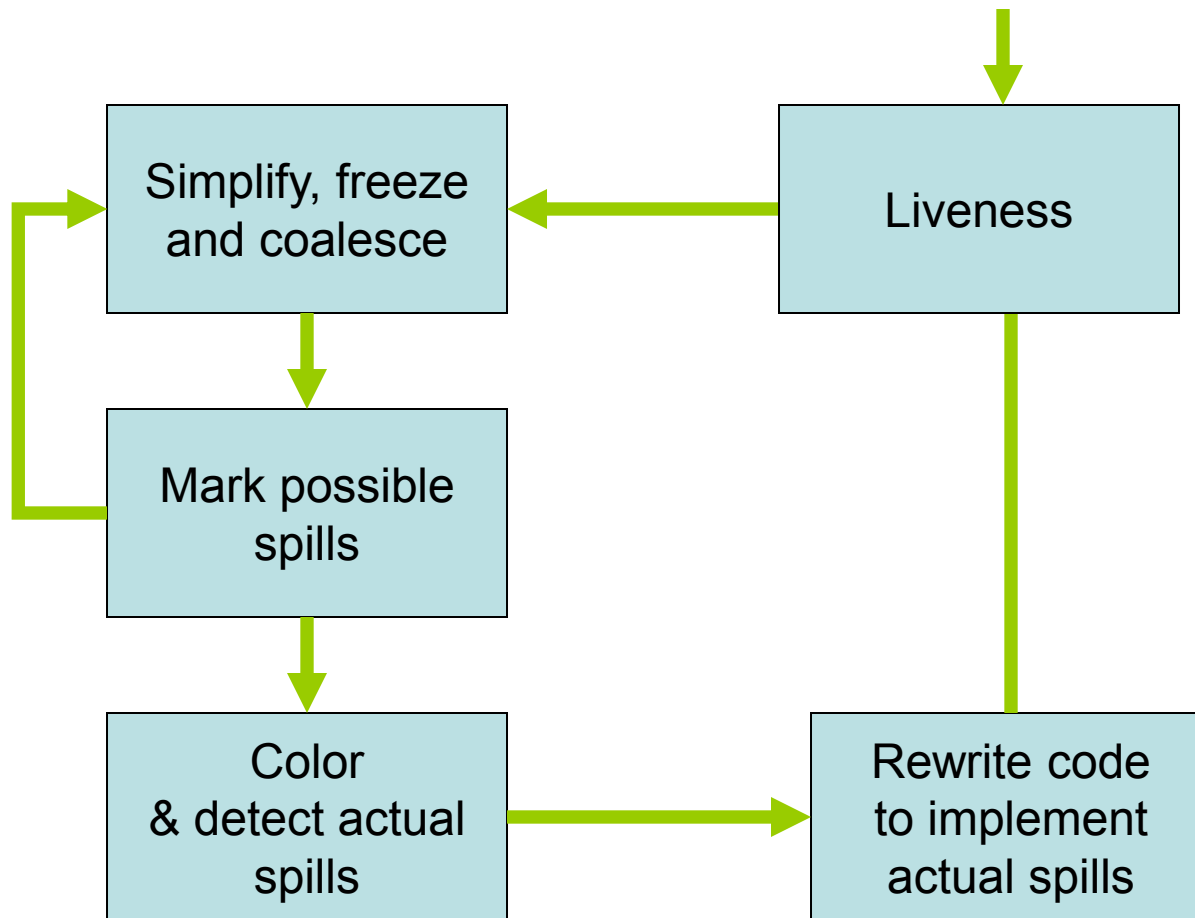


- Solution 1 (Briggs): avoid creation of high-degree ($\geq K$) nodes
- Solution 2 (George): a can be coalesced with b if every neighbour t of a :
 - already interferes with b , or
 - has low-degree ($< K$)

Simplify & Coalesce

- **Step 1 (simplify):** simplify as much as possible without removing nodes that are the source or destination of a move (*move-related nodes*)
- **Step 2 (coalesce):** coalesce move-related nodes provided low-degree node results
- **Step 3 (freeze):** if neither steps 1 or 2 apply, freeze a move instruction: registers involved are marked *not move-related* and try step 1 again

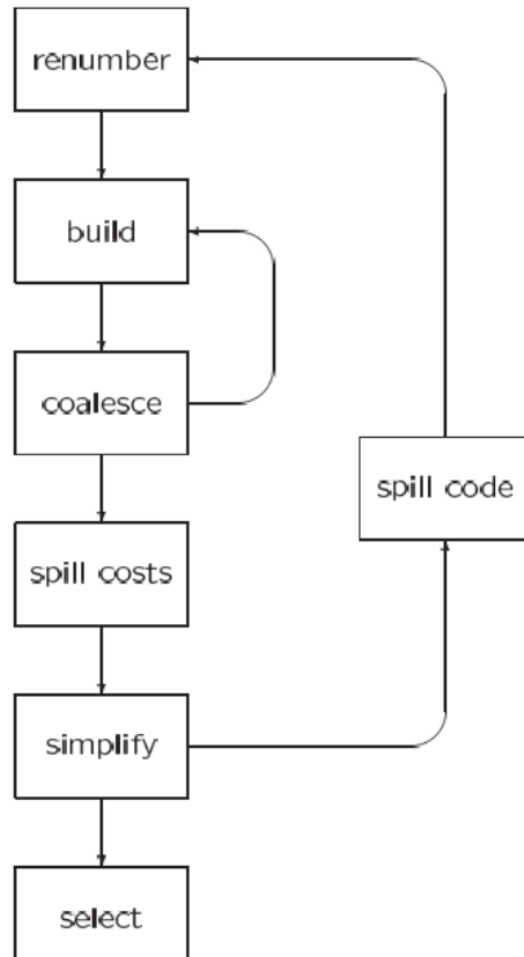
Overall Algorithm



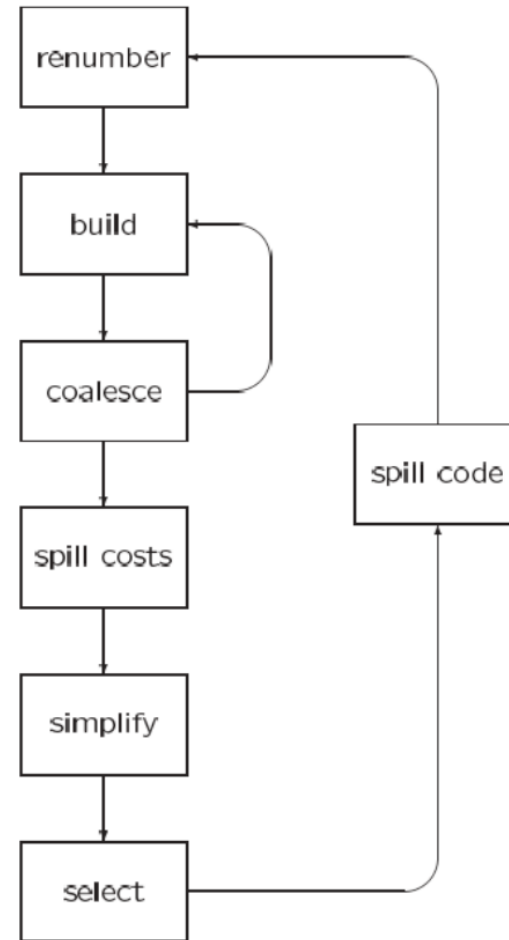
Variations

May read papers for more variations

Chaitin



Briggs



Questions?

Linear scan

- Given the live ranges of variables in a function, the algorithm scans all the live ranges in a **single pass**, allocating registers to variables in a **greedy fashion**.
- M. Poletto, V. Sarkar. *Linear scan register allocation*. 1999.

LINEARSCANREGISTERALLOCATION

$active \leftarrow \{\}$

foreach live interval i , in order of increasing start point

 EXPIREOLDINTERVALS(i)

if length($active$) = R **then**

 SPILLATINTERVAL(i)

else

$register[i] \leftarrow$ a register removed from pool of free registers
 add i to $active$, sorted by increasing end point


```
EXPIREOLDINTERVALS(i)  
  foreach interval j in active, in order of increasing end point  
    if  $endpoint[j] \geq startpoint[i]$  then  
      return  
  remove j from active  
  add register[j] to pool of free registers
```

SPILLATINTERVAL(i)

$spill \leftarrow$ last interval in $active$

if $endpoint[spill] > endpoint[i]$ **then**

$register[i] \leftarrow register[spill]$

$location[spill] \leftarrow$ new stack location


remove $spill$ from $active$

add i to $active$, sorted by increasing end point

else

$location[i] \leftarrow$ new stack location

Example




- Initially, *active* is empty



of available registers is $R = 2$


Example



- Step 1:
 $active = \langle A \rangle$

of available registers is $R = 2$


Example



- Step 2:
 $active = \langle A, B \rangle$

of available registers is $R = 2$


Example



- Step 3:
- 3 live intervals overlap

of available registers is $R = 2$


Example



of available registers is $R = 2$

- Step 3:
- 3 live intervals overlap
- spills C , whose interval ends furthest away from the current point
- $active = \langle A, B \rangle$


Example



- Step 4:
- A expires
- $active = \langle A, B, D \rangle$

of available registers is $R = 2$


Example



- Step 5:
- B expires
- $active = \langle \overline{B}, D, E \rangle$

of available registers is $R = 2$


Example



- In the end, C is the only variable not allocated to a register.

of available registers is $R = 2$

Example



of available registers is $R = 2$

- In the end, C is the only variable not allocated to a register.
- Otherwise, both one of A and B and one of D and E would have been spilled to memory.

Questions?

Conclusion

- # of slides
 - Graph coloring: ~50
 - Linear scan: ~10

Conclusion

- # of slides
 - Graph coloring: ~50
 - Linear scan: ~10
- Linear scan is much more simpler!
 - Only about 10% slower than a perfectly implemented graph coloring algorithm
 - And your code may not be that perfect 😊

Acknowledgements

- Graph coloring slides are adapted from *Register Allocation* by David Walker
- Linear scan pseudo-code and example are adapted from *Linear Scan Register Allocation* by M. Poletto and V. Sarkar.